

# [8.5] Signals

-by Sandesh Paudel


- > operating system uses **exceptions** to support a form of **exceptional control flow** known as **context switch**.
- >**linux signal**: a higher-level software form of exceptional control flow that allows processes and the kernel to interrupt other processes.
- > **signal**: a small message that notifies a process that an **event** of some type has occurred in the system. Textbook 757 has 30 different types of signals that are supported on Linux.

<i><b>ID</b></i>	<i><b>Name</b></i>	<i><b>Default Action</b></i>	<i><b>Corresponding Event</b></i>
2	SIGINT	Terminate	User typed ctrl-c
9	SIGKILL	Terminate	Kill program (cannot override or ignore)
11	SIGSEGV	Terminate	Segmentation violation
14	SIGALRM	Terminate	Timer signal
17	SIGCHLD	Ignore	Child stopped or terminated

- > low-level hardware exceptions are processed by the **kernel's exception handlers** and are not normally visible to user processes. **Signals** provide a mechanism for exposing the occurrence of **such exceptions** to user processes.
- For eg, if a process attempts to divide by zero, then the kernel sends a **SIGFPE** signal. Other signals correspond to other such events.
- > If you type Ctrl+C while a process is running in the foreground, then the kernel sends a **SIGINT** to each process in the foreground process group. A process can forcibly terminate another process by sending it a **SIGKILL** signal.

***Sending a signal.*** The kernel *sends (delivers)* a signal to a destination process by updating some state in the context of the destination process. The signal is delivered for one of two reasons: (1) The kernel has detected a system event such as a divide-by-zero error or the termination of a child process. (2) A process has invoked the `kill` function (discussed in the next section) to explicitly request the kernel to send a signal to the destination process. A process can send a signal to itself.

***Receiving a signal.*** A destination process *receives* a signal when it is forced by the kernel to react in some way to the delivery of the signal. The process can either ignore the signal, terminate, or *catch* the signal by executing a user-level function called a *signal handler*. Figure 8.27 shows the basic idea of a handler catching a signal.



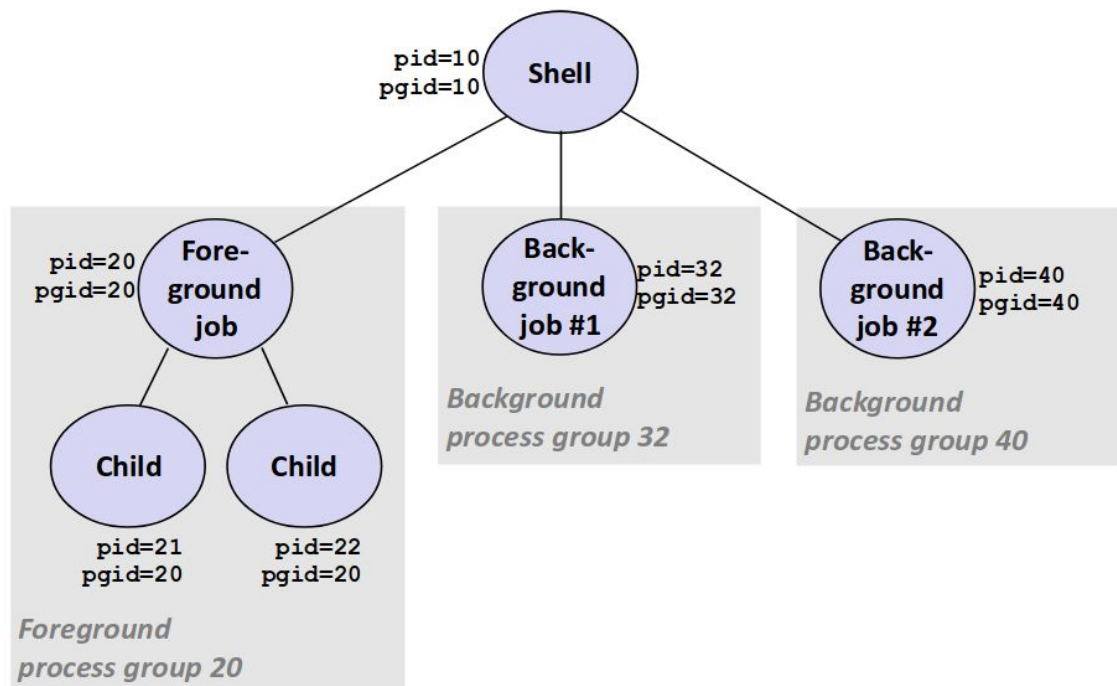
> **Pending Signal:** a signal that has been sent but not yet received. There can be at most one pending signal of a particular type at a time.

- If a process has a pending signal of type  $k$ , then any subsequent signals of type  $k$  sent to that process are not queued; they are **discarded**.
- A process can selectively **block** the receipt of certain signals. The blocked signal will be delivered but not received until the process unblocks it.
- For each process, the kernel maintains a set of pending signals in the **pending bit vector**, and a set of blocked signals in the **blocked bit vector**.
- The kernel sets bit  $k$  in pending whenever a signal of type  $k$  is delivered and clears bit  $k$  in pending whenever a signal of type  $k$  is received.

#### > Sending Signals

- **Process groups:** every process belongs to exactly one process group identified by a positive integer **process group ID**. It can be accessed with `getpgrp()`.

- Every process belongs to exactly one process group



```
#include <unistd.h>
```

```
pid_t getpgrp(void);
```

Returns: process group ID of calling process

- By default, a child process belongs to the same process group as its parent. A process can change the process group of itself or another process by using the `setpgid` function [from pid to pgid]:

```
#include <unistd.h>
```

```
int setpgid(pid_t pid, pid_t pgid);
```

Returns: 0 on success, -1 on error

## Sending Signals with the kill Function

Processes send signals to other processes (including themselves) by calling the `kill` function.

```
#include <sys/types.h>
```

```
#include <signal.h>
```

```
int kill(pid_t pid, int sig);
```

Returns: 0 if OK, -1 on error

If `pid` is greater than zero, then the `kill` function sends signal number `sig` to process `pid`. If `pid` is equal to zero, then `kill` sends signal `sig` to every process in the process group of the calling process, including the calling process itself. If `pid` is less than zero, then `kill` sends signal `sig` to every process in process group `|pid|` (the absolute value of `pid`). Figure 8.29 shows an example of a parent that uses the `kill` function to send a `SIGKILL` signal to its child.

## Sending Signals with the /bin/kill Program

The /bin/kill program sends an arbitrary signal to another process. For example, the command

```
linux> /bin/kill -9 15213
```

sends signal 9 (SIGKILL) to process 15213. A negative PID causes the signal to be sent to every process in process group PID. For example, the command

```
linux> /bin/kill -9 -15213
```

sends a SIGKILL signal to every process in process group 15213. Note that we use the complete path /bin/kill here because some Unix shells have their own built-in kill command.

## Sending Signals with the alarm Function

A process can send SIGALRM signals to itself by calling the alarm function.

```
#include <unistd.h>

unsigned int alarm(unsigned int secs);
                Returns: remaining seconds of previous alarm, or 0 if no previous alarm
```

### > Sending Signals via Keyboard:

- **Job:** abstraction to represent the **processes** that are created as a result of evaluating single command line. There can be at most one foreground job but multiple or zero background jobs.
- **Pipe:** output of one process acts as the input of another. `linux> ls | sort`
-

# Receiving Signals

- Kernel is returning from an exception handler and is ready to pass control to process p
- Kernel computes the set of pending & nonblocked signals for process p (PNB set)
- If (PNB is empty)
  - Pass control to next instruction in the logical flow for p
- Else
  - Choose least nonzero bit k in pnb and force process p to **receive** signal k
  - The receipt of the signal triggers some **action** by p
  - Repeat for all nonzero k in pnb
  - Pass control to next instruction in logical flow for p

# Default Actions

- Each signal type has a predefined **default action**, which is one of:
  - The process terminates
  - The process stops until restarted by a SIGCONT signal
  - The process ignores the signal



# Installing Signal Handlers

- The signal function modifies the default action associated with the receipt of signal `signum`:
  - `handler_t *signal(int signum, handler_t *handler)`
- Different values for handler:
  - `SIG_IGN`: ignore signals of type `signum`
  - `SIG_DFL`: revert to the default action on receipt of signals of type `signum`
  - Otherwise, `handler` is the address of a user-level function (signal handler)
    - Called when process receives signal of type `signum`
    - Referred to as “installing” the handler
    - Executing handler is called “catching” or “handling” the signal
    - When the handler executes its return statement, control passes back to instruction in the control flow of the process that was interrupted by receipt of the signal

## Signal Handling Example

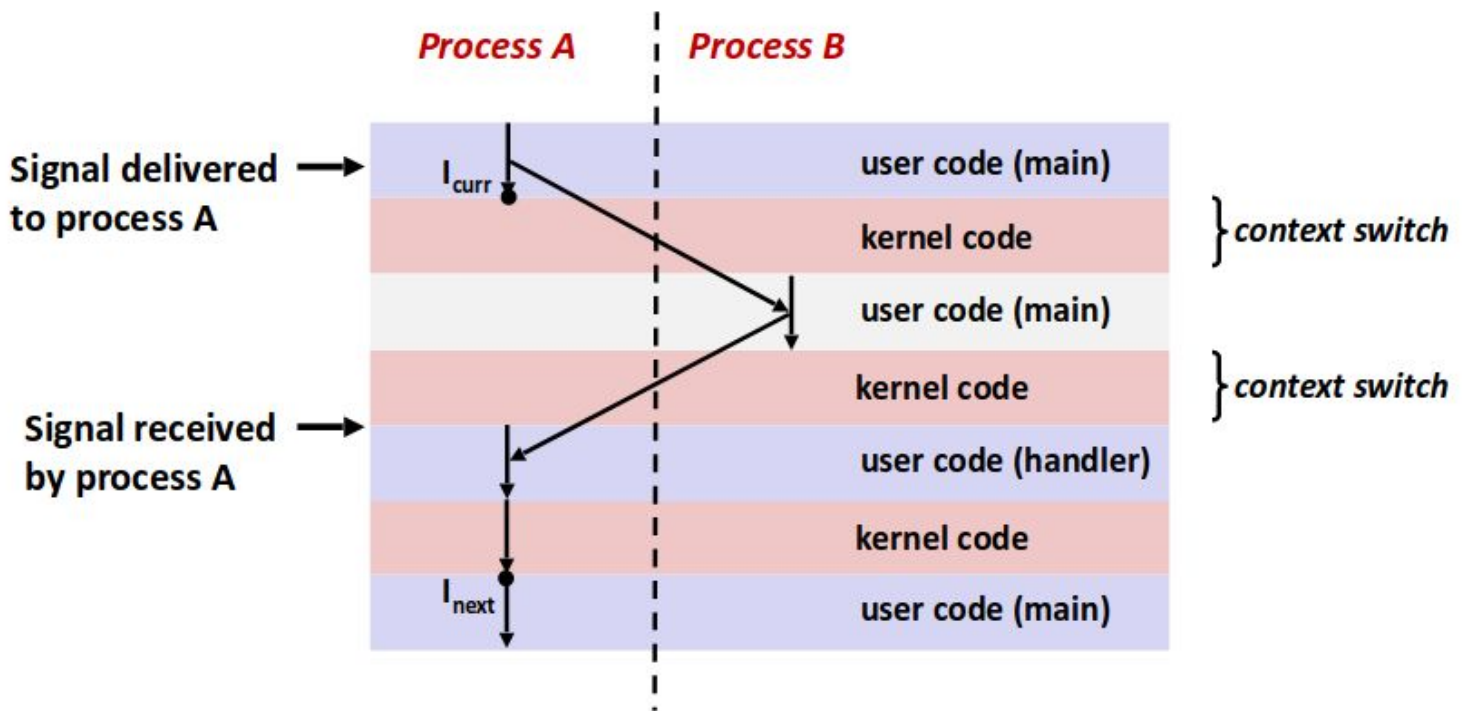
```
void sigint_handler(int sig) /* SIGINT handler */
{
    printf("So you think you can stop the bomb with ctrl-c, do you?\n");
    sleep(2);
    printf("Well...\n");
    fflush(stdout);
    sleep(1);
    printf("OK. :-)\n");
    exit(0);
}

int main()
{
    /* Install the SIGINT handler */
    if (signal(SIGINT, sigint_handler) == SIG_ERR)
        unix_error("signal error");

    /* Wait for the receipt of a signal */
    pause();

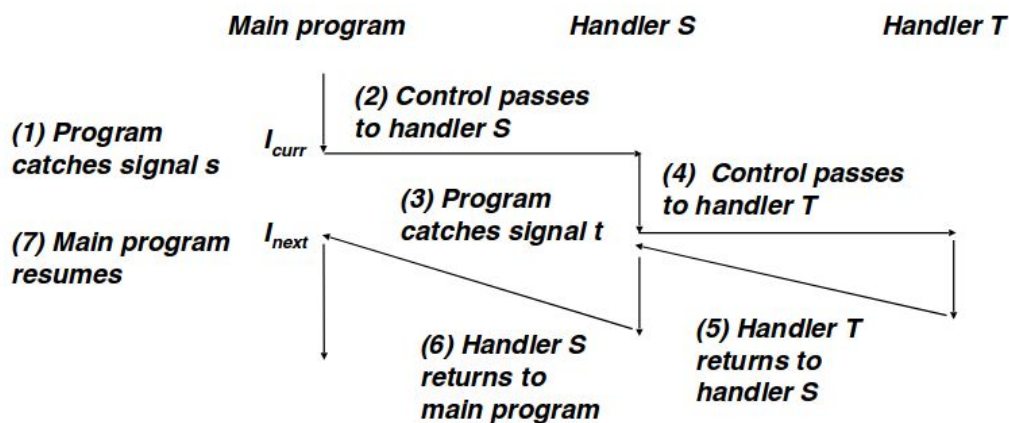
    return 0;
}
```

sigint.c



## Nested Signal Handlers

- Handlers can be interrupted by other handlers





# Blocking Signals

```
sigset_t mask, prev_mask;

sigemptyset(&mask);
sigaddset(&mask, SIGINT);

/* Block SIGINT and save previous blocked set */
sigprocmask(SIG_BLOCK, &mask, &prev_mask);

/* Code region that will not be interrupted by SIGINT */

/* Restore previous blocked set, unblocking SIGINT */
sigprocmask(SIG_SETMASK, &prev_mask, NULL);
```

- Explicit blocking and unblocking signal
  - sigprocmask function
  - sigemptyset – Create empty set
  - sigfillset – Add every signal number to set
  - sigaddset – Add signal number to set
  - sigdelset – Delete signal number from set

# Safe Signal Handling

- Handlers are tricky because they are **concurrent with main program** and may **share the same global data structures**.

```
static int x = 5;
void handler(int sig)
{
    x = 10;
}

int main(int argc, char **argv)
{
    int pid;
    Signal(SIGCHLD, handler);

    if ((pid = Fork()) == 0) { /* Child */
        Execve("/bin/date", argv, NULL);
    }

    if (x == 5)
        y = x * 2; // You'd expect y == 10
    exit(0);
}
```

What if the following happens:

- Parent process executes and finishes **if (x == 5)**
- Context switch to child, which then terminates, sends a SIGCHLD signal
- Another context switch back to parent, and now the kernel needs to execute the SIGCHLD handler
- When return to parent process, **y == 20!**

3

# Safe Signal Handling

- Handlers are tricky because they are **concurrent with main program** and may **share the same global data structures**.
  - Programmers have no control over the execution ordering between the main program and the signal handler, that is:
    - when a signal happens/delivers (depends on user or other process)
    - when the signal handler will be executed (depends on kernel)
  - If not careful, shared data structures can be corrupted



# Fixing the Signal Handling Bug

```
static int x = 5;
void handler(int sig)
{
    x = 10;
}

int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, prev_all;
    sigfillset(&mask_all);
    signal(SIGCHLD, handler);

    if ((pid = Fork()) == 0) { /* Child */
        Execve("/bin/date", argv, NULL);
    }

    Sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
    if (x == 5)
        y = x * 2; // You'd expect y == 10
    Sigprocmask(SIG_SETMASK, &prev_all, NULL);

    exit(0);
}
```

- Block all signals before accessing a shared, global data structure.
- Can't use a lock (later in this course)

34

## Async-Signal-Safety

- Function is **async-signal-safe** if either reentrant (e.g., no access to global variables) or non-interruptible by signals.
- Posix guarantees 117 functions to be async-signal-safe
  - Source: “man 7 signal”
  - Popular functions on the list:
    - `_exit`, `write`, `wait`, `waitpid`, `sleep`, `kill`
  - Popular functions that are **not** on the list:
    - `printf`, `sprintf`, `malloc`, `exit`
    - Unfortunate fact: `write` is the only async-signal-safe output function

# Another Unsafe Signal Handler Example

- Assume a program wants to do the following:
  - The parent creates multiple child processes
  - When each child process is created, add the child PID to a queue
  - When a child process terminates, the parent process removes the child PID from the queue
- One possible implementation:
  - An array for keeping the child PIDs
  - Use a loop to fork child, and add PID to the array after fork
  - Install a handler for SIGCHLD in parent process
  - The SIGCHLD handler removes the child PID

## First Attempt

```
void handler(int sig)
{
    pid_t pid;

    while ((pid = wait(NULL)) > 0) { /* Reap child */
        deletejob(pid); /* Delete the child from the job list */
    }
}

int main(int argc, char **argv)
{
    int pid;

    signal(SIGCHLD, handler);
    initjobs(); /* Initialize the job list */

    while (1) {
        if ((pid = Fork()) == 0) { /* Child */
            Execve("/bin/date", argv, NULL);
        }
        addjob(pid); /* Add the child to the job list */
    }
    exit(0);
}
```

---



## Second Attempt

```
void handler(int sig)
{
    sigset_t mask_all, prev_all;
    pid_t pid;

    sigfillset(&mask_all);
    while ((pid = wait(NULL)) > 0) {
        sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
        deletejob(pid);
        sigprocmask(SIG_SETMASK, &prev_all, NULL);
    }
}

int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, prev_all;

    sigfillset(&mask_all);
    signal(SIGCHLD, handler);
    initjobs(); /* Initialize the job list */

    while (1) {
        if ((pid = Fork()) == 0) {
            Execve("/bin/date", argv, NULL);
        }
        sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
        addjob(pid);
        sigprocmask(SIG_SETMASK, &prev_all, NULL);
    }
    exit(0);
}
```

The following can happen:

- Child runs, and terminates
- Kernel sends SIGCHLD
- Context switch to parent, but before it can run, kernel has to handle SIGCHLD first
- The handler deletes the job, which does nothing
- The parent process resumes and adds a terminated child to job list

38

## Third Attempt (The Correct One)

```
int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, mask_one, prev_one;

    Sigfillset(&mask_all);
    Sigemptyset(&mask_one);
    Sigaddset(&mask_one, SIGCHLD);
    signal(SIGCHLD, handler);
    initjobs(); /* Initialize the job list */

    while (1) {
        Sigprocmask(SIG_BLOCK, &mask_one, &prev_one); /* Block SIGCHLD */
        if ((pid = Fork()) == 0) { /* Child process */
            Sigprocmask(SIG_SETMASK, &prev_one, NULL); /* Unblock SIGCHLD */
            Execve("/bin/date", argv, NULL);
        }
        Sigprocmask(SIG_BLOCK, &mask_all, NULL); /* Parent process */
        addjob(pid); /* Add the child to the job list */
        Sigprocmask(SIG_SETMASK, &prev_one, NULL); /* Unblock SIGCHLD */
    }
    exit(0);
}
```

39