

THREADS

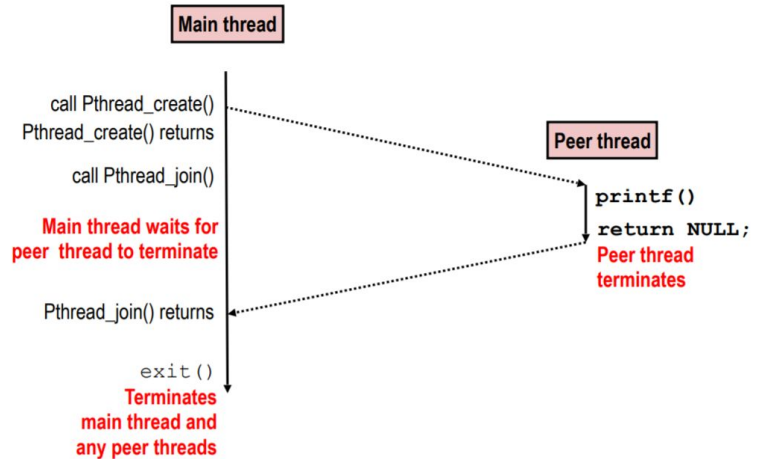
By Sandesh Paudel

Posix Threads (Pthreads) Interface

- **Pthreads**: Standard interface for ~60 functions that manipulate threads from C programs

- Creating and reaping threads
 - `pthread_create()`
 - `pthread_join()`
- Determining your thread ID
 - `pthread_self()`
- Terminating threads
 - `pthread_cancel()`
 - `pthread_exit()`
 - `exit()` [terminates all threads], `RET` [terminates current thread]
- Synchronizing access to shared variables
 - `pthread_mutex_init`
 - `pthread_mutex_unlock`

Execution of Threaded “hello, world”



- **Key observation**: In general, any sequentially consistent interleaving is possible, but some give an unexpected result!

i (thread) instr_i %rdx₁ %rdx₂ cnt

1	L ₁	0	-	0
1	U ₁	1	-	0
1	S ₁	1	-	1
2	L ₂	-	1	1
2	U ₂	-	2	1
2	S ₂	-	2	2

Thread 1 critical section

Thread 2 critical section

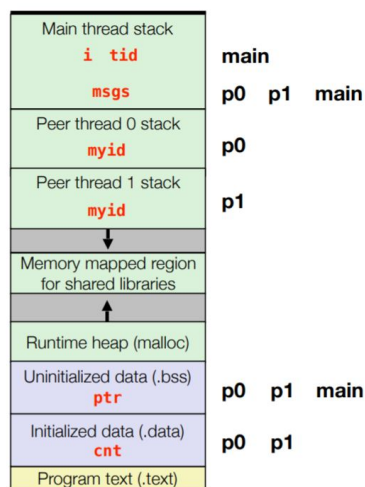
Example Program to Illustrate Sharing

```

char **ptr; /* global var */
void *thread(void *vargp)
{
    long myid = (long)vargp;
    static int cnt = 0;
    printf("[%ld]: %s (cnt=%d)\n",
           myid, ptr[myid], ++cnt);
    return NULL;
}

int main()
{
    long i;
    pthread_t tid;
    char *msgs[2] = {
        "Hello from foo",
        "Hello from bar"
    };
    ptr = msgs;
    for (i = 0; i < 2; i++)
        pthread_create(&tid,
                       NULL,
                       thread,
                       (void *)i);
    pthread_exit(NULL);
}
  
```

sharing.c



i (thread) instr_i %rdx₁ %rdx₂ cnt

1	L ₁	0	-	0
1	U ₁	1	-	0
2	L ₂	-	0	0
1	S ₁	1	-	1
2	U ₂	-	1	1
2	S ₂	-	1	1

Assembly Code for Counter Loop

C code for counter loop in thread i


```
for (i = 0; i < niters; i++)
    cnt++;
```

Asm code for thread i

```
movq (%rdi), %rcx
testq %rcx, %rcx
jle .L2
movl $0, %eax
.L3:
movq cnt(%rip), %rdx
addq $1, %rdx
movq %rdx, cnt(%rip)
addq $1, %rax
cmpq %rcx, %rax
jne .L3
.L2:
```

Annotations:

- H_i : Head
- L_i : Load cnt
- U_i : Update cnt
- S_i : Store cnt
- T_i : Tail



Def: A trajectory is *safe* iff it does not enter any unsafe region

Claim: A trajectory is *correct* (wrt cnt) iff it is safe

Terminology:

- **Binary semaphore:** semaphore whose value is always 0 or 1
- **Mutex:** binary semaphore used for mutual exclusion
 - P operation: “locking” the mutex
 - V operation: “unlocking” or “releasing” the mutex
 - “Holding” a mutex: locked and not yet unlocked.
- **Counting semaphore:** used as a counter for set of available resources.

Semaphores = Locks

Proper Synchronization

- Define and initialize a mutex for the shared variable `cnt`:

```
volatile long cnt = 0; /* Counter */
sem_t mutex;          /* Semaphore that protects cnt */
Sem_init(&mutex, 0, 1); /* mutex = 1 */
```


- Surround critical section with P and V:

```
for (i = 0; i < niters; i++) {
    P(&mutex);
    cnt++;
    V(&mutex);
}
```

goodcnt.c

```
linux> ./goodcnt 10000
OK cnt=20000
linux> ./goodcnt 10000
OK cnt=20000
linux>
```

Warning: It's orders of magnitude slower than `badcnt.c`.



Removing an item from a shared buffer:

```
/* Remove and return the first item from buffer sp */
int sbuf_remove(sbuf_t *sp)
{
    int item;

    P(&sp->mutex);          /* Lock the buffer */
    item = sp->buf[(++sp->front)%(sp->n)]; /* Remove the item */
    V(&sp->mutex);          /* Unlock the buffer */

    return item;
}
```

Inserting an item into a shared buffer:

```
/* Insert item onto the rear of shared buffer sp */
void sbuf_insert(sbuf_t *sp, int item)
{
    P(&sp->mutex);          /* Lock the buffer */
    sp->buf[(++sp->rear)%(sp->n)] = item; /* Insert the item */
    V(&sp->mutex);          /* Unlock the buffer */
}
```

Implementation

Removing an item from a shared buffer:

```
/* Remove and return the first item from buffer sp */
int sbuf_remove(sbuf_t *sp)
{
    int item;
    P(&sp->items);           /* Wait for available item */
    P(&sp->mutex);          /* Lock the buffer */
    item = sp->buf[(++sp->front)%(sp->n)]; /* Remove the item */
    V(&sp->mutex);          /* Unlock the buffer */
    V(&sp->slots);          /* Announce available slot */
    return item;
}
```

Inserting an item into a shared buffer:

```
/* Insert item onto the rear of shared buffer sp */
void sbuf_insert(sbuf_t *sp, int item)
{
    P(&sp->slots);           /* Wait for available slot */
    P(&sp->mutex);          /* Lock the buffer */
    sp->buf[(++sp->rear)%(sp->n)] = item; /* Insert the item */
    V(&sp->mutex);          /* Unlock the buffer */
    V(&sp->items);          /* Announce available item */
}
```

Deadlock

- Def: A process/thread is *deadlocked* if and only if it is waiting for a condition that will never be true
- General to concurrent/parallel programming (threads, processes)
- Typical Scenario
 - Processes 1 and 2 needs two resources (A and B) to proceed
 - Process 1 acquires A, waits for B
 - Process 2 acquires B, waits for A
 - Both will wait forever!

Tid[0]:	Tid[1]:
P(s ₀);	P(s ₁);
P(s ₁);	P(s ₀);
cnt++;	cnt++;
V(s ₀);	V(s ₁);
V(s ₁);	V(s ₀);



Tid[0]:	Tid[1]:
P(s ₀);	P(s ₀);
P(s ₁);	P(s ₁);
cnt++;	cnt++;
V(s ₀);	V(s ₁);
V(s ₁);	V(s ₀);

How About Using a Mutex?

```
static int x = 5;
void handler(int sig)
{
    P(&mutex);
    x = 10;
    V(&mutex);
}

int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, prev_all;
    signal(SIGCHLD, handler);

    if ((pid = Fork()) == 0) { /* Child */
        Execve("/bin/date", argv, NULL);
    }

    P(&mutex);
    if (x == 5)
        y = x * 2; // You'd expect y == 10
    V(&mutex);

    exit(0);
}
```

- This implementation will get into a deadlock.
- Signal handler wants the mutex, which is acquired by the main program.
- **Key**: signal handler is in the same process as the main program. The kernel forces the handler to finish before returning to the main program.


Amdahl's Law

- Gene Amdahl (1922 – 2015). Giant in computer architecture
- Captures the difficulty of using parallelism to speed things up
- Amdahl's Law
 - f: Parallelizable fraction of a program
 - N: Number of processors (i.e., maximal achievable speedup)

$$\text{Speedup} = \frac{1}{1 - f + \frac{f}{N}}$$

- Completely parallelizable (f = 1): Speedup = N
- Completely sequential (f = 0): Speedup = 1


Why the Sequential Bottleneck?



- Maximum speedup limited by the sequential portion
- Main cause: **Non-parallelizable operations on data**
- Parallel portion is usually not perfectly parallel as well
 - e.g., Synchronization overhead

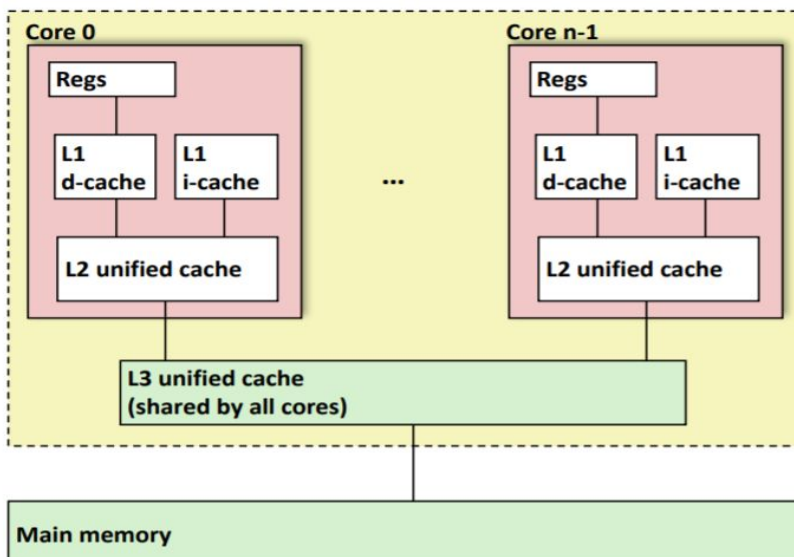
Any benefits?

- Can single-core multi-threading provide any performance gains?
- If Thread A has a cache miss and the pipeline gets stalled, switch to Thread C. Improves the overall performance.



Multi-core for multi threading

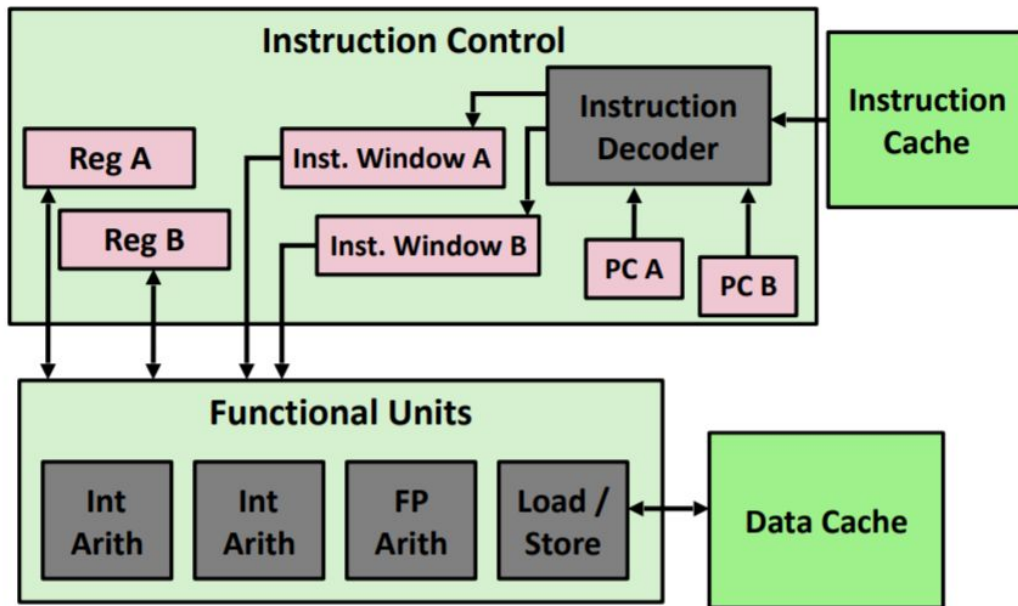
Typical Multi-core Processor



- Traditional multiprocessing: symmetric multiprocessor (SMP)
- Every core is exactly the same. Private registers, L1/L2 caches, etc.
- Share L3 (LLC) and main memory

Hyper-threading

- Intel's terminology. More commonly known as: Simultaneous Multi-threading (SMT)
- Replicate enough hardware structures to process K instruction streams
- K copies of all registers. Share functional units





Conventional Multi-threading vs. Hyper-threading

Conventional
Multi-threading

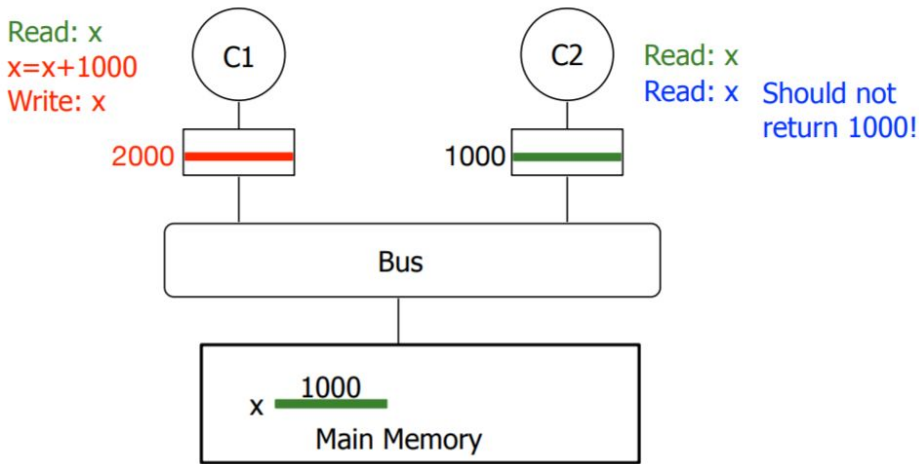
Hyper-threading

- Thread 1
- Context Switch
- Thread 2



Multiple threads actually execute in parallel (even with one single core)

No/little context switch overhead




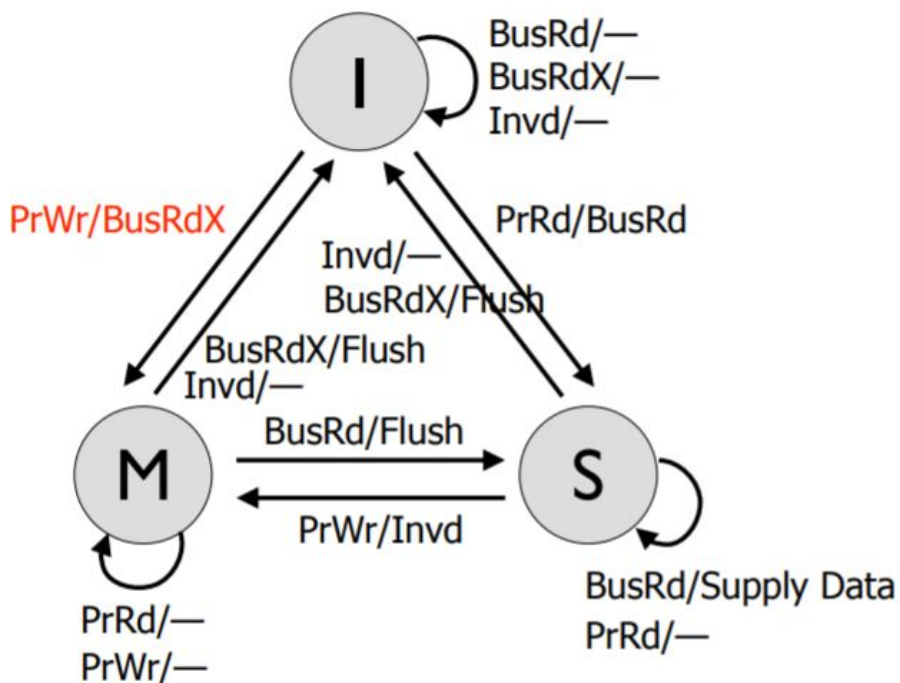
35

Invalidate-Based Cache Coherence

Associate each cache line with 3 states: **Modified, Invalid, Shared**

Below: State Transition for x in C2's cache;
Syntax: Event/Action

Write: x = 7000



38

Does Hardware Have to Keep Cache Coherent?

- Hardware-guaranteed cache coherence is complex to implement.
- Can the programmers ensure cache coherence themselves?
- Key: ISA must provide cache flush/invalidate instructions
 - FLUSH-LOCAL A: Flushes/invalidates the cache block containing address A from a processor's local cache.
 - FLUSH-GLOBAL A: Flushes/invalidates the cache block containing address A from all other processors' caches.
 - FLUSH-CACHE X: Flushes/invalidates all blocks in cache X.
- Classic example: TLB
 - Hardware does not guarantee that TLBs of different core are coherent
 - ISA provides instructions for OS to flush PTEs
 - Called "TLB shutdown"

Take CSC 251/ECE 204 to learn more about advanced computer architecture concepts.